

Postgres Indexes

Zach Croft

Postgres Internals Resource: <https://www.interdb.jp/pg/index.html>

Index: What, Why?


```
[postgres=# SELECT relname, oid, relfilenode FROM pg_class WHERE relname = 'idx_grades_id_covering';
   relname    |   oid   | relfilenode
-----+-----+-----
 idx_grades_id_covering | 24594 |      24594
(1 row)
```

```
[root@9c50520968ae:/# ls -la $PGDATA/base/* | grep 24594
-rw----- 1 postgres postgres 112336896 Jun 12 01:32 24594
```

- Data structure to enhance database performance.
 - analogous to using binary search instead of linear search on a sorted list, or categorizing library books by sections but not one-by-one.
- Supported index types: B-tree, Hash, GiST, SP-GiST, GIN, BRIN, and bloom.
 - PG primary and secondary keys use indexes and default to B-tree.

Heap and Pages

```
typedef struct PageHeaderData {src/include/storage/bufpage.h
{
    /* XXX LSN is member of *any* block, not only page-organized ones */
    PageLogRecPtr pd_lsn;           /* LSN: next byte after last byte of xlog
                                     * record for last change to this page */
    uint16 pd_checksum;             /* checksum */
    uint16 pd_flags;               /* flag bits, see below */
    LocationIndex pd_lower;         /* offset to start of free space */
    LocationIndex pd_upper;         /* offset to end of free space */
    uint16 pd_special;             /* offset to start of special space */
    uint16 pd_pagesize_version;
    TransactionId pd_prune_xid;    /* oldest prunable XID, or zero if none */
    ItemIdData pd_lnp[1];           /* beginning of line pointer array */
} PageHeaderData;
typedef PageHeaderData *PageHeader;
typedef uint64 XLogRecPtr;
```




- Data files (table, index, etc.) are divided into pages (or blocks) of a fixed default length of 8KB, each identified with a block number.
- Tables are organized in a ‘heap’ structure (this is distinct from the heap tree structure AND the heap related to memory allocation).
- Heap tables contain ‘pages’ (this is also distinct from the page related to memory allocation, those are typically 4KB) with header metadata, line pointer(s), free space, and heap tuple(s). A DB page could span 3 OS pages.
 - Line pointers grow ‘up’ and heap tuples grow ‘down’ similar to heap and stack in memory.
- Heap tuples store the actual records and are identified by a **tuple identifier (TID)** built from the block number and the offset number (line pointer ID).
 - Line pointers are essentially mini-indexes within a page. Don’t confuse this with a PostgreSQL index.


Note: these images only show the **key**.

B-tree (Original)

- B-trees (bee trees) are balanced trees with pages as nodes.
 - Each node contains elements, which are key-value pairs from key to TID.
 - Postgres does not actually map to the primary key, which is not universal between different DB engines.
 - If a node has m child nodes, then it can have at most $m-1$ elements.




- (left) I added sequentially, (right) added in a different order.
 - Ordering of the tree is dependent on input order.



B-tree Limitations

Storage

- All nodes (root, internal, leaf) have key-value pair elements (we will see, not necessary), taking more storage, therefore taller trees, and therefore more IO to disk (fetches more OS pages).
 - Indexes may not entirely fit in memory if they are large—the database might have to use the swap which is slow.




Performance: Range Query

- Consider the above b-tree on id:
 - SELECT id FROM grades WHERE id ≥ 1 AND id < 8 ;
 - To get to id=1, fetches 3 pages.
 - id=2 will probably be cached unless moved into swap.
 - id=4 is probably cached.
 - Id=5, fetches 2 pages.
 - id=6 is probably cached.
 - We got id=8 whether we wanted it or not—page IO fetches everything in it.

B+tree (Fancy)

Consequent nodes are doubly-linked which makes range queries much easier.

- B+trees are balanced trees with pages as nodes, but:
 - Internal nodes (everything except leaf nodes), are just keys.
 - Nodes are pages; height is reduced and therefore traversal times are reduced.
- Some keys are duplicated, and searches always go to the leaf nodes.
- Postgres often operates by keeping all internal nodes in memory and then leaving the leaf nodes in the heap for quicker traversal.



Postgres really uses B+trees (modified) in all instances that there is a B-tree. B+trees are not always better as there are trade-offs, and also other indexes might be a better fit for situations.

```

[postgres=# \d grades
              Table "public.grades"
  Column |      Type       | Collation | Nullable |
          | character varying(255) |           | not null |
  id    | integer        |           | not null |
  grade | integer        |           | not null |
  name  | character varying(255) |           | not null |
Indexes:
  "grades_pkey" PRIMARY KEY, btree (id)
[postgres=# SELECT COUNT(*) FROM grades;
 count
-----
 5000000
(1 row)

```

How Do Indexes Perform?

```

Seq Scan on grades (cost=0.00..102028.00 rows=2970027
  Filter: ((grade >= 0) AND (grade < 60))
  Rows Removed by Filter: 2026244
Planning Time: 0.315 ms
JIT:
  Functions: 4
  Options: Inlining false, Optimization false, Expressions
  Timing: Generation 0.504 ms (Deform 0.137 ms), Inlining
  Execution Time: 334.864 ms
(9 rows)
Time: 337.576 ms

```

```

Indexes:
  "idx_grades_id_covering" btree (id) INCLUDE (grade)

```

```

postgres-# EXPLAIN ANALYZE
postgres-# SELECT id
postgres-# ,grade
postgres-# FROM grades
postgres-# WHERE id < 1500 AND grade > 75;
-----
Index Only Scan using idx_grades_id_coverin
  Index Cond: (id < 1500)
  Filter: (grade > 75)
  Rows Removed by Filter: 1144

```




Table Scan

- Generally, the slowest way queries can be made.
 - DB engines often use multiple workers for speed.
 - Buffer system and cache layers help.

Index Scan

- A typical index scan where a key-value pair is found and then the corresponding heap table page is fetched.
- `SELECT *` and `SELECT name` are same (row-based).

Index-Only Scan

- Never jump to the heap!
- All requested values are in the index file.
- Composite indexes (including other values in your index) work well here.

Bitmap Index Scan

- Works with bitmap heap scans to create a bitmap for heap pages to include.
- `BitmapAnd` takes the intersection of multiple bitmaps.

Shopify



“In one high-throughput system at Shopify we’ve seen a 50 percent decrease in INSERT statement duration by switching from UUIDv4 to ULID for idempotency keys.”

Line Pointers (mini-indexes) And Clustering

If I don't really mind how long insertion takes, can I just sort the heap table pages and heap tuples on insertion and skip the IO's from the index? [Index-organized tables...](#) are not in Postgres. But you can cluster periodically.

Demo!

Final Thoughts And Discussion

- Indexes are generally useful but can slow down smaller queries, the planner takes note of this and may skip indexes.
- Indexes help most on massive tables, but horizontal partitioning and/or sharding can help by splitting massive tables into smaller tables.
- How do indexes work with locking and transactions?

```
rails_zcroft=# \di
```

```
public | y  
(836 rows)
```